

HBase

FEATURES OF NOSQL DATABASES

Generic data model

- Heterogeneous containers, including sets, maps, and arrays

Dynamic type discovery and conversion

- NoSQL analytics systems support runtime type identification and conversion so that custom business logic can be used to

dictate analytic treatment of variation.

Non-relational and De-normalised

- Data is stored in single tables as compared to joining multiple tables.

Commodity hardware

- Adding more of the economical servers allows NoSQL databases to scale to handle more data.

Highly distributable

- Distributed databases can store and process a set of information on more than one device.

WHAT IS HBASE?

Apache HBase is the Hadoop database, a distributed, column oriented, scalable, big data store.

- Use Apache HBase when you need random, realtime read/write access to your Big Data.
- This project's goal is the hosting of very large tables - billions of rows, X millions of columns - atop clusters of commodity

hardware.

- Apache HBase is an open-source, distributed, versioned, non-relational database modelled after Google's Big table: A

Distributed Storage System for Structured Data by Chang et al.

- Just as Big table leverages the distributed data storage provided by the Google File System, Apache HBase provides Big

table-like capabilities on top of Hadoop and HDFS.

HBASE DATA MODEL — COLUMNS & COLUMN FAMILIES

```
{  
  "1" : {  
    "A" : {  
      "foo": "x"  
    },  
    "aaaaa" : {  
      "A" : {  
        "foo" : "y"  
      },  
      "aaaab" : {  
        "A" : {  
          "foo" : "world"  
        },  

```

1. Top Level Key/map pair is called row
2. A is called as Column Family
3. Foo is the column of family A

HBASE DATA MODEL – LOGICAL REPRESENTATION

Customer id		Customer Address data			Customer order data		
RowKey		CF1			CF2		
		colA	colB	colC	colA	colB	colC
axxx		Val		val	val		val
gxxx			val			val	

Data is **accessed and stored together**:

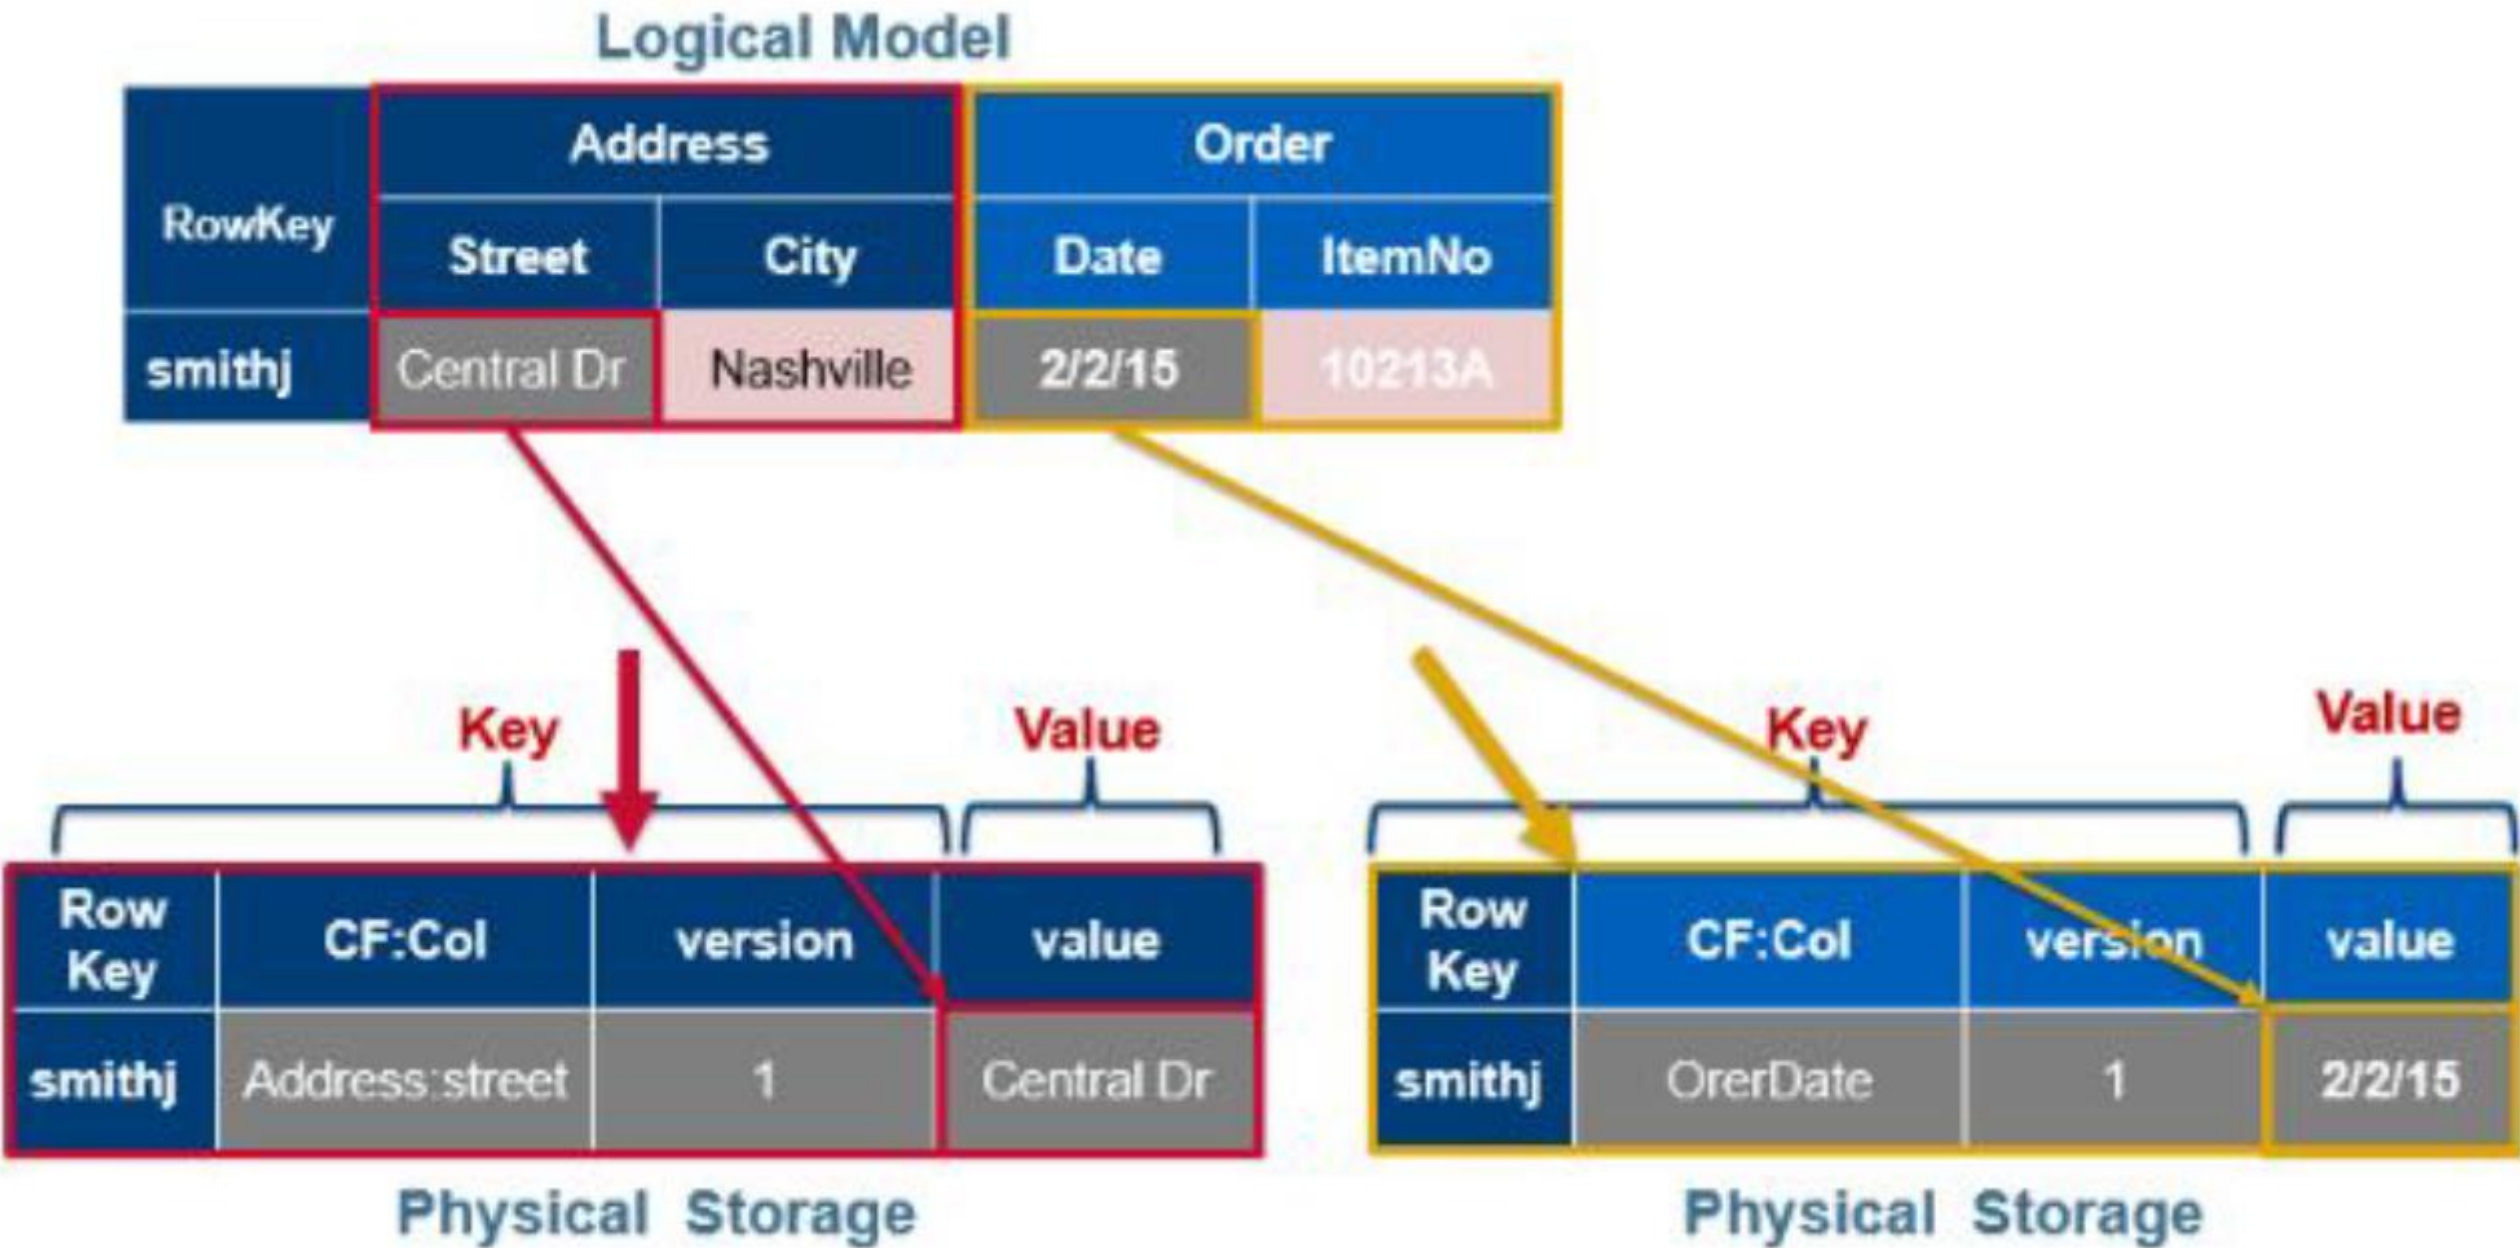
- RowKey is the primary index
- Column Families group similar data by **row key**

Cell Coordinates= **Key**

Value

Row key	Column Family	Column Qualifier	Timestamp	Value
Smithj	Address	city	1391813876369	Nashville

LOGICAL VERSUS PHYSICAL REPRESENTATION



CAP Theorem

Consistency - This means that the data in the database remains consistent after the execution of an operation. For

example after an update operation, all clients see the same data.

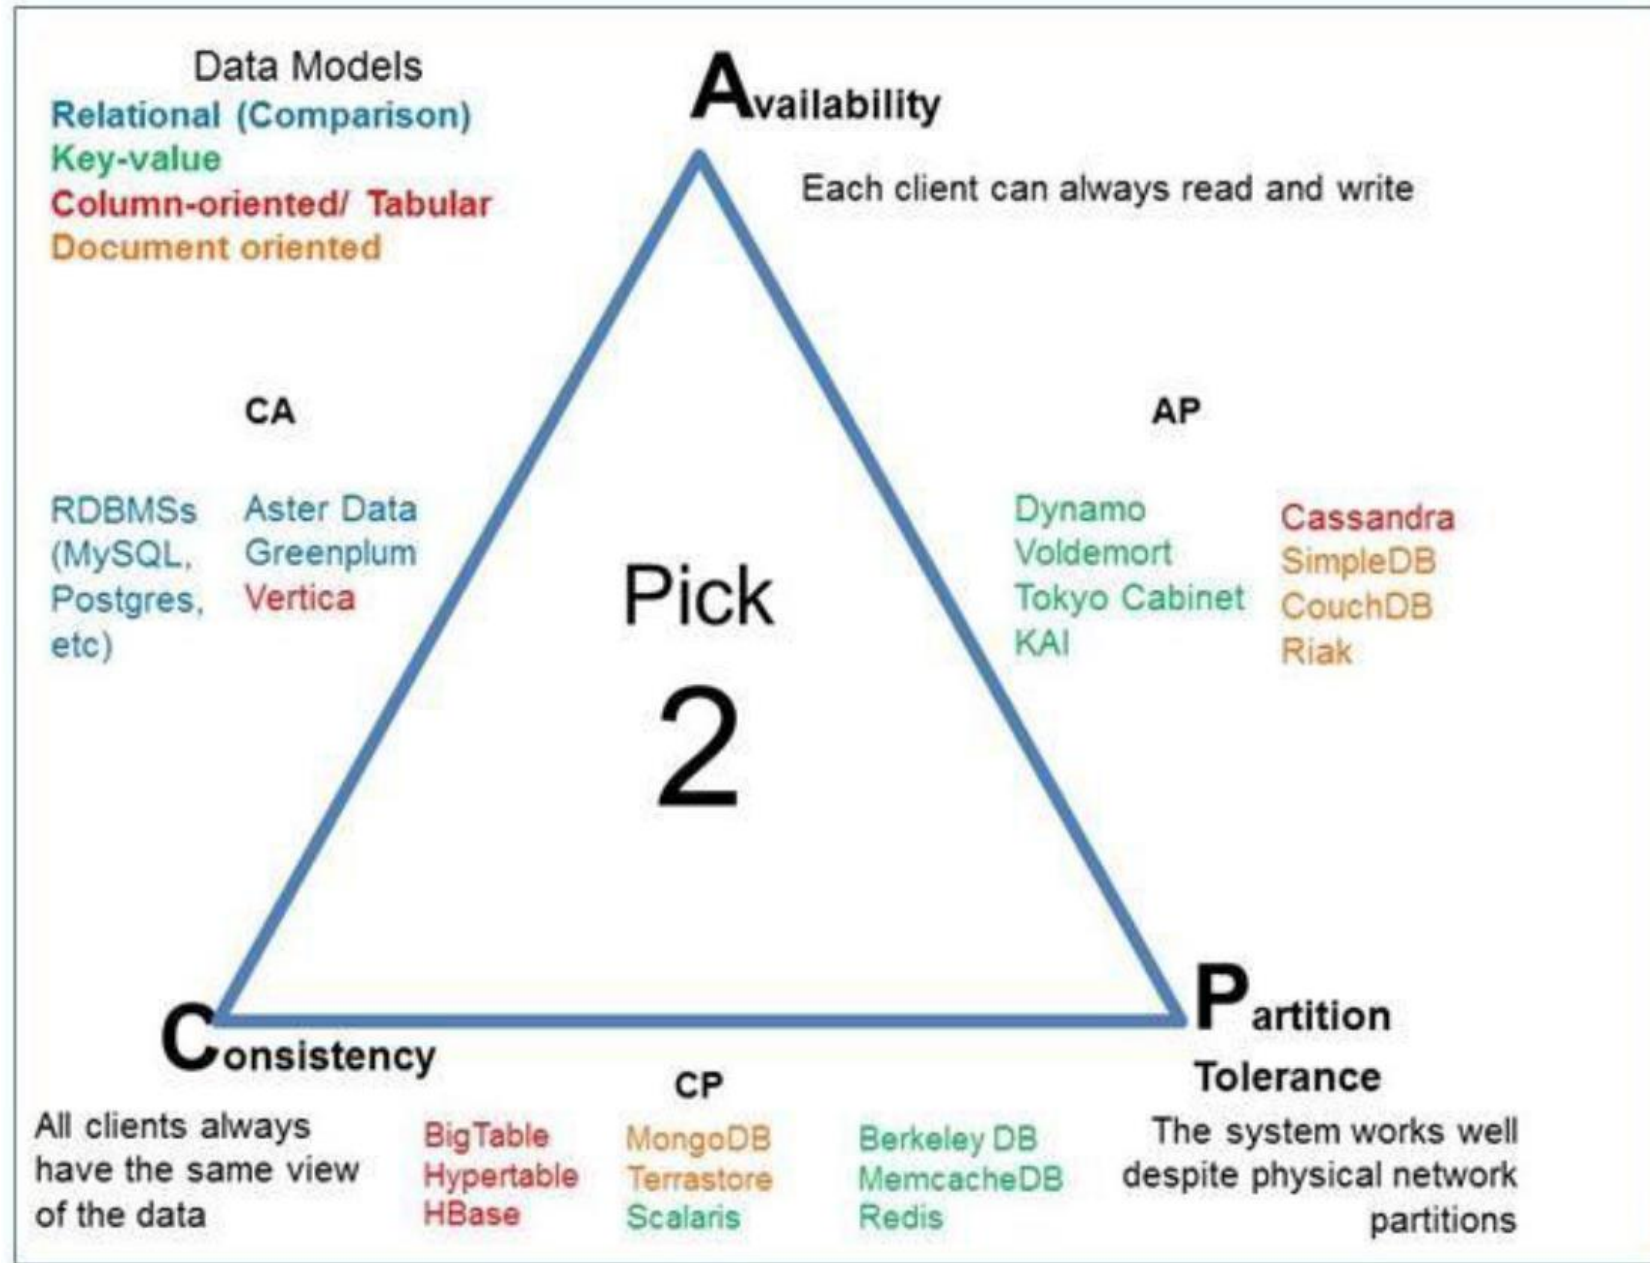
- **Availability** - This means that the system is always on (service guarantee availability), no downtime.

- **Partition Tolerance** - This means that the system continues to function even if the communication among the servers is unreliable, i.e. the servers may be partitioned into multiple groups that cannot communicate with one another.

Duplicate Copy of same data is maintained on multiple machines.

- This increases availability, but decreases consistency.

CAP



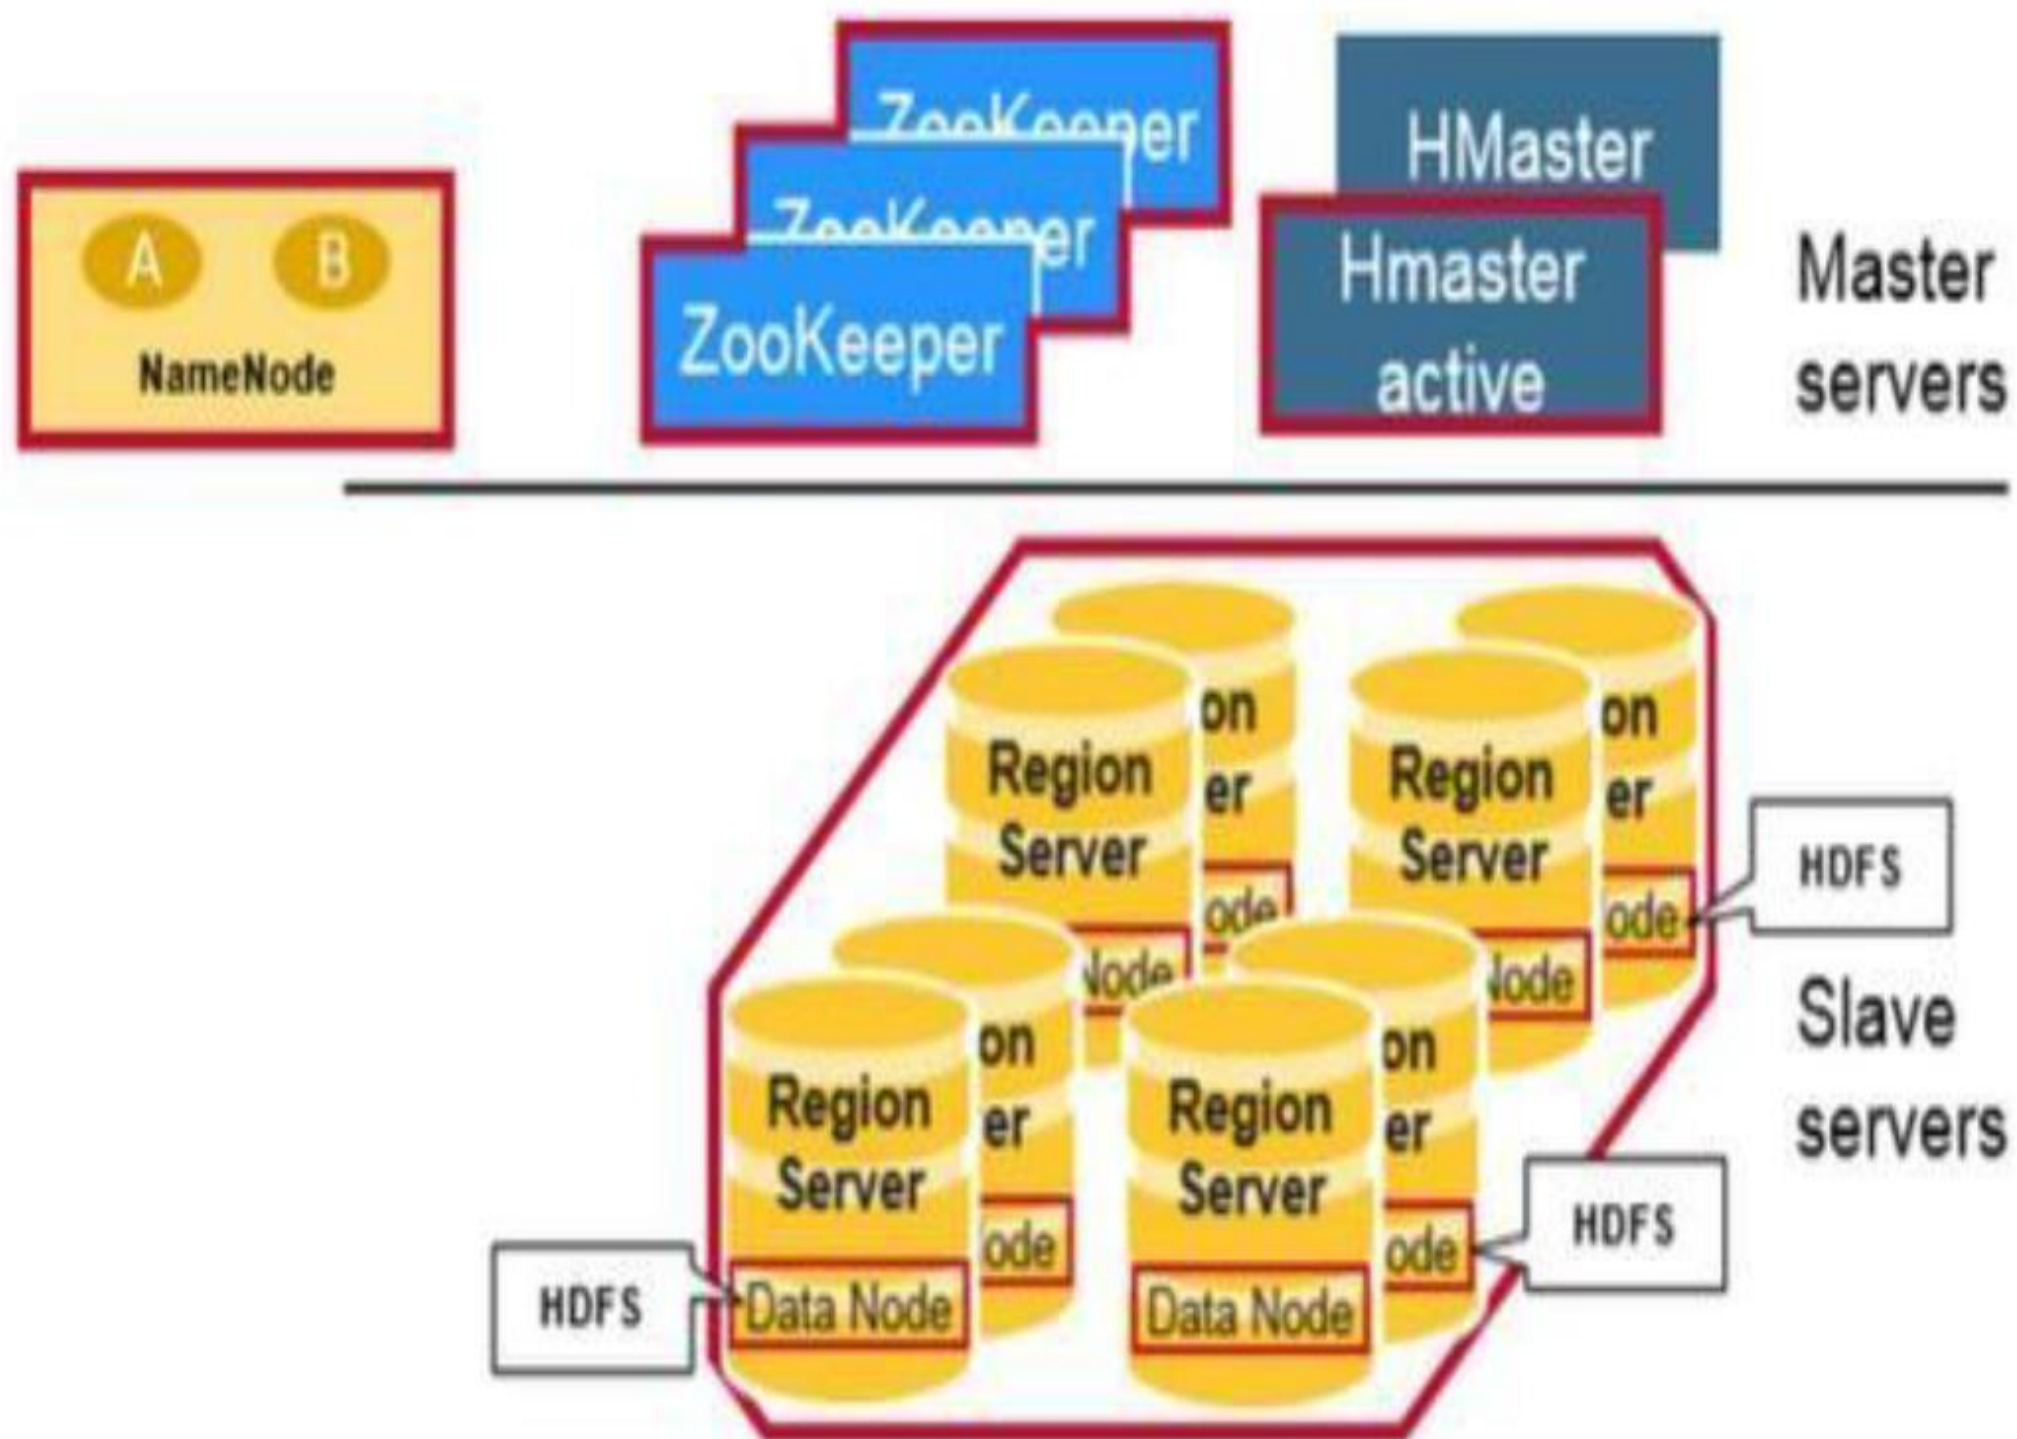
Hbase Architecture

HBase is composed of three types of servers in a master slave type of architecture.

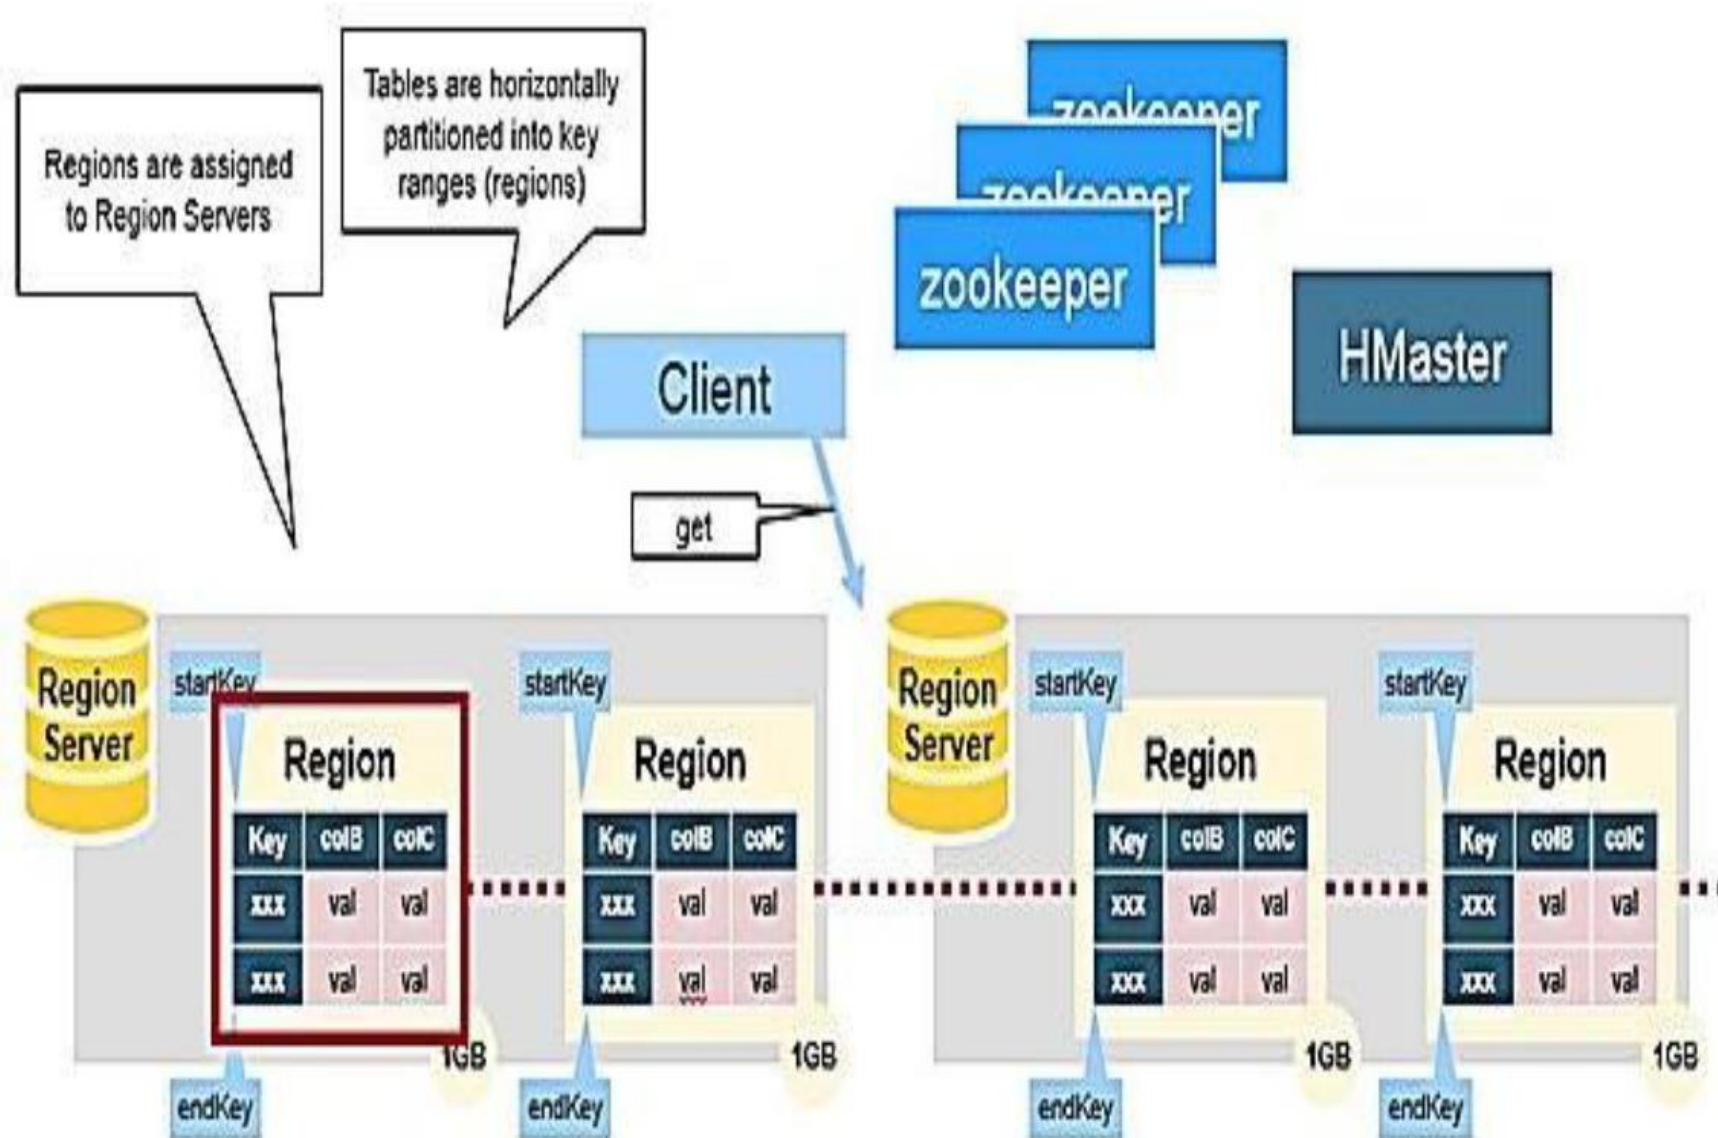
- Region servers serve data for reads and writes
- HBase Master process handles the Region assignment, DDL (create, delete tables) operations
- Zookeeper maintains a live cluster state

The Hadoop DataNode stores the data that the Region Server is managing

- All HBase data is stored in HDFS files
- The NameNode maintains metadata information for all the physical data blocks that comprise the files



REGIONS



HMASTER

Region assignment, DDL (create, delete tables) operations are handled by the HBase Master.

A master is responsible for:

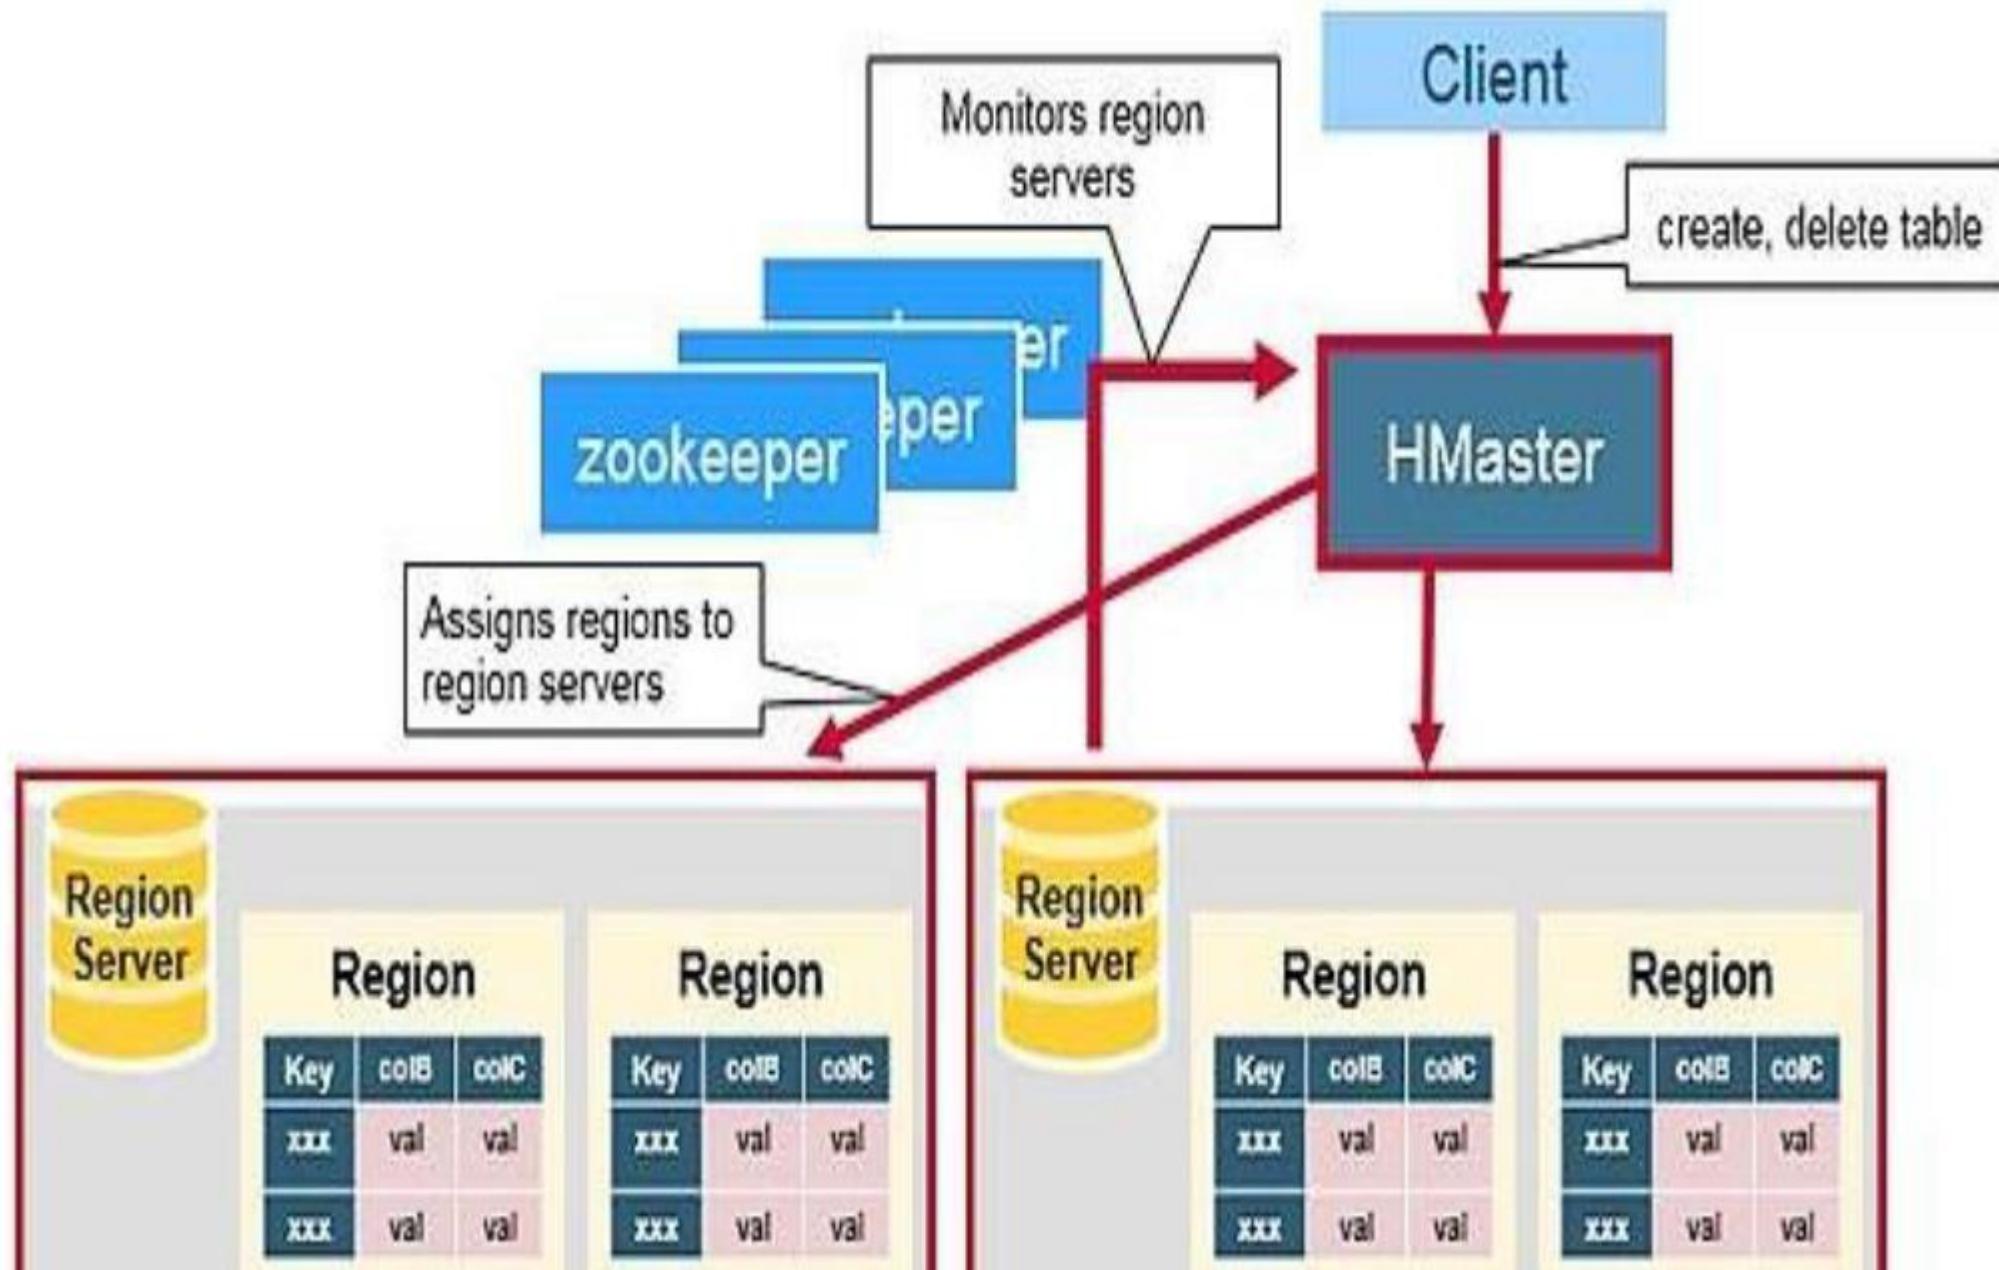
- Coordinating the region servers
- Assigning regions on startup
- Re-assigning regions for recovery or load balancing
- Monitoring all RegionServer instances in the cluster

(listens for notifications from
zookeeper)

Admin functions

- Interface for creating, deleting, updating tables

HMASTER



HBASE FIRST READ

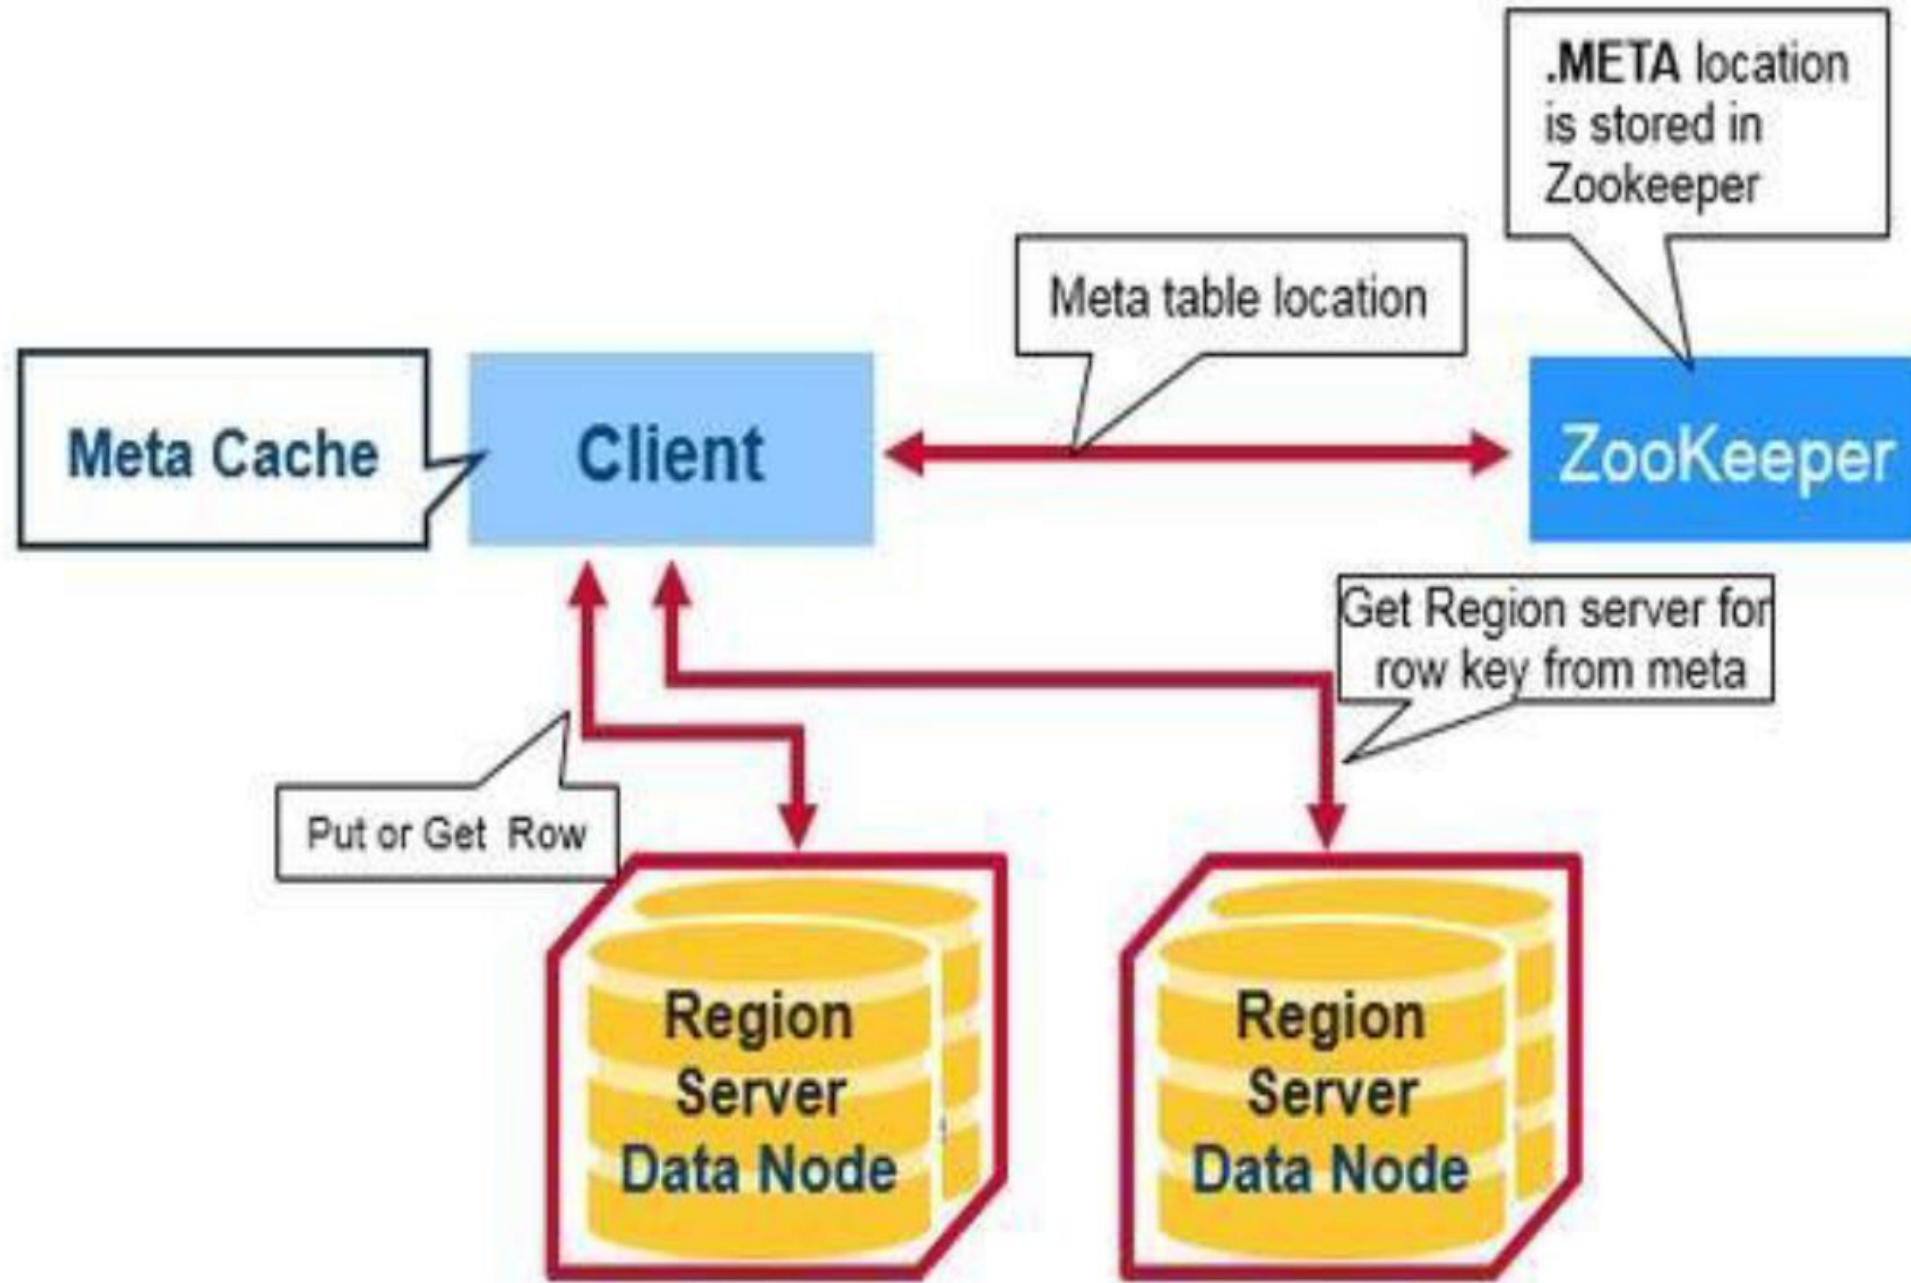
There is a special HBase Catalog table called the META table, which holds the location of the regions in the cluster

- ZooKeeper stores the location of the META table
- The client gets the Region server that hosts the META table from ZooKeeper
- The client will query the .META. server to get the region server corresponding to the row key it wants to access. The client

caches this information along with the META table location

- It will get the row from the corresponding Region Server
- For future reads, the client uses the cache to retrieve the META location and previously read row keys
- Over time, it does not need to query the META table, unless there is a miss because a region has moved; then it will re-query and update the cache

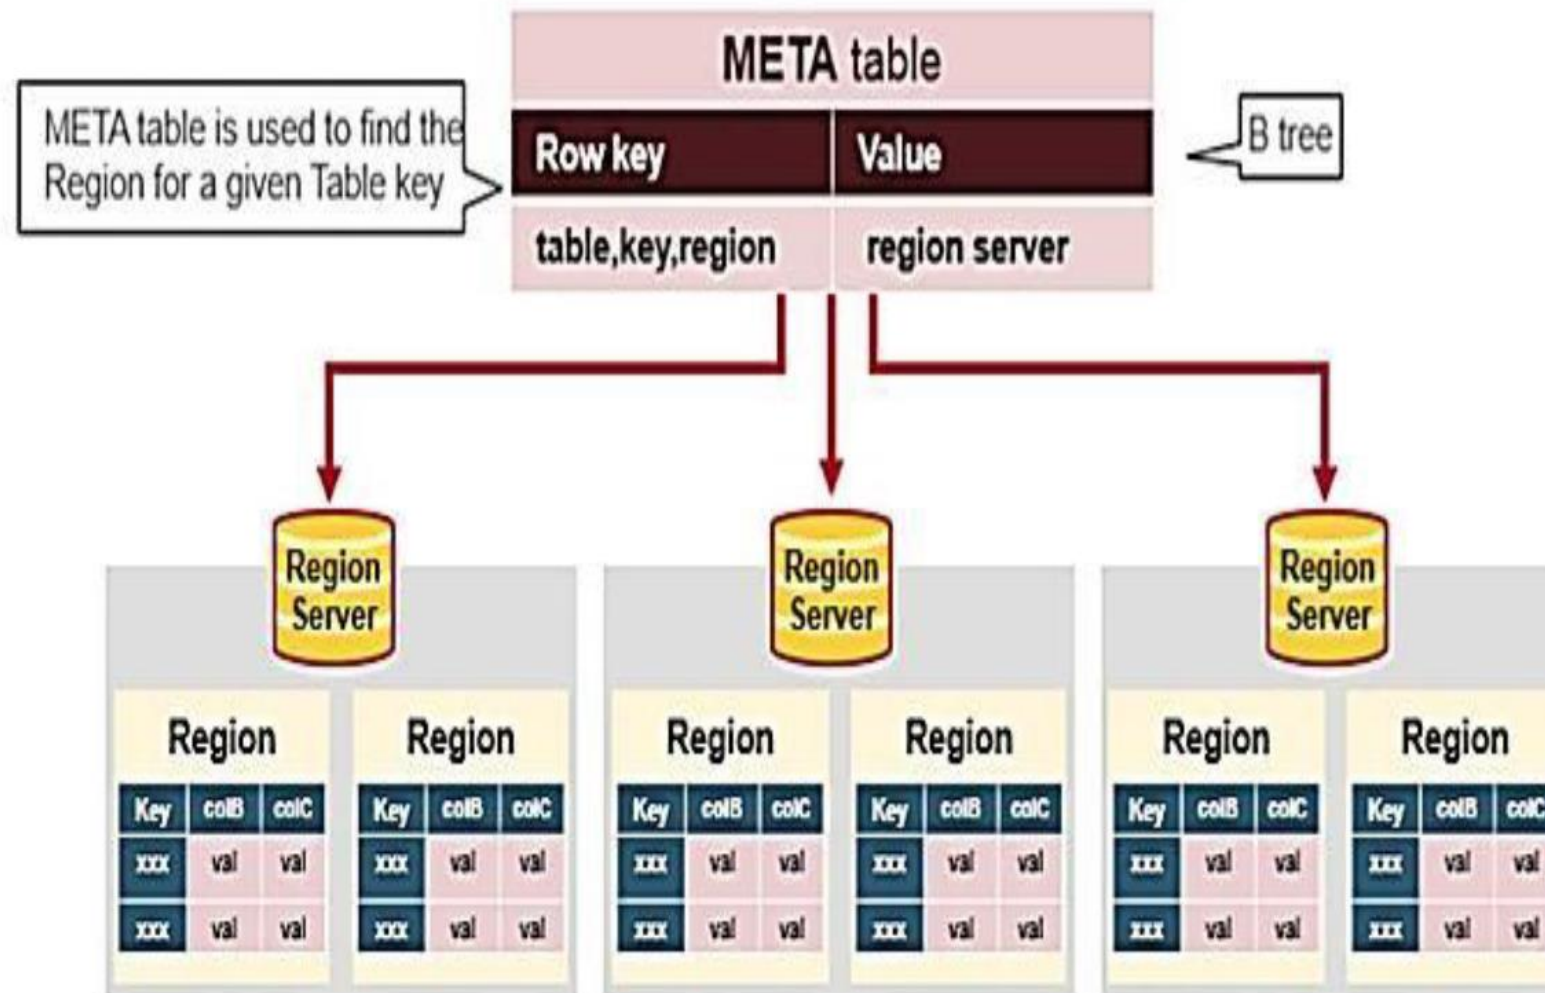
HBASE FIRST READ



HBASE META TABLE

This META table is an HBase table that keeps a list of all regions in the system

- The .META. table is like a b tree
- The .META. table structure is as follows: Key: region start key, region id Values: RegionServer



REGION SERVER COMPONENTS

Region Server runs on an HDFS data node and has the following components:

WAL

- Write Ahead Log is a file on the distributed file system. The WAL is used to store new data that hasn't yet been persisted to permanent storage; it is used for recovery in the case of failure.

BlockCache

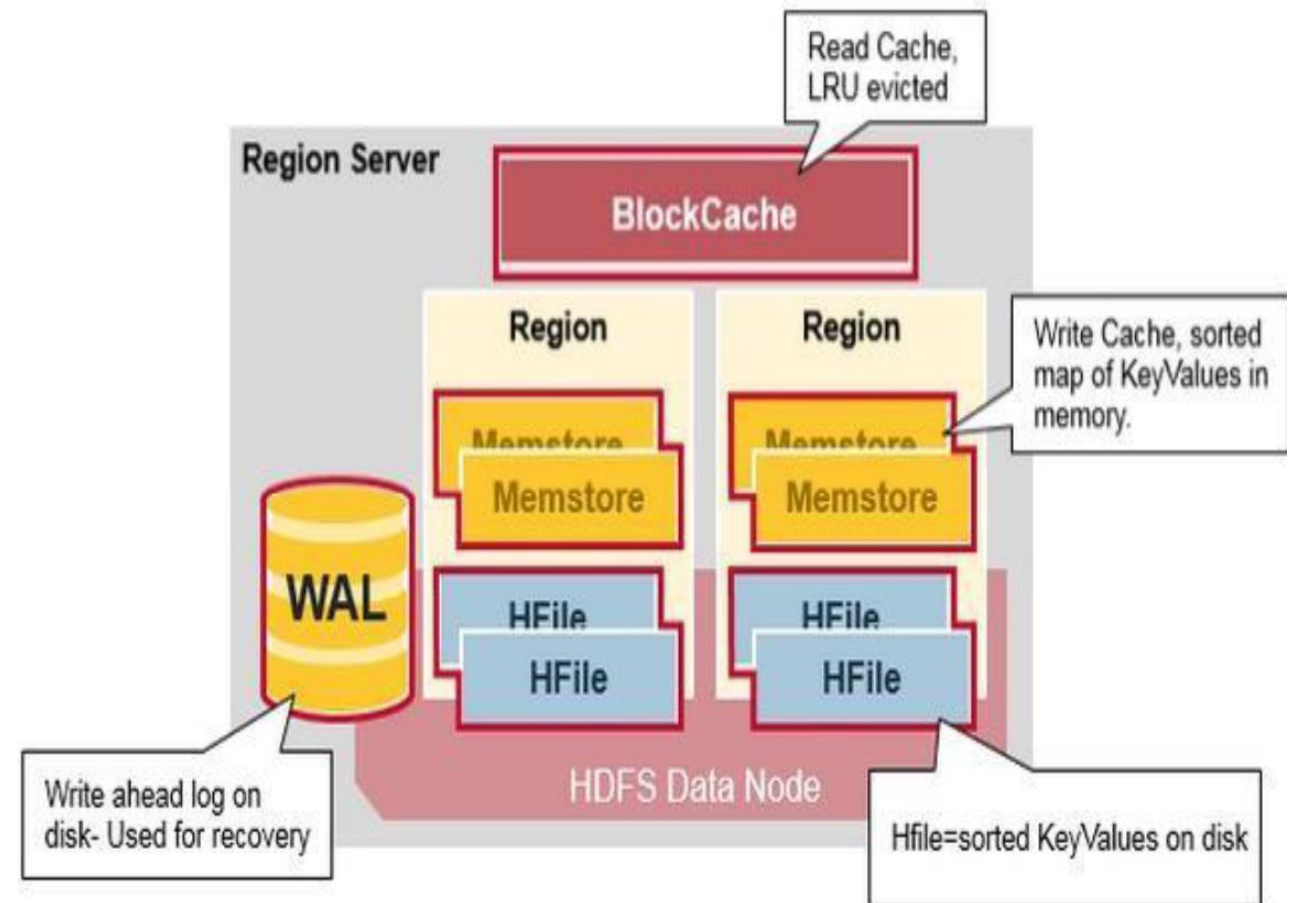
- It is the read cache. It stores frequently read data in memory. Least Recently Used data is evicted when full.

MemStore

- It is the write cache. It stores new data which has not yet been written to disk. It is sorted before writing to disk. There is one MemStore per column family per region.

Hfiles

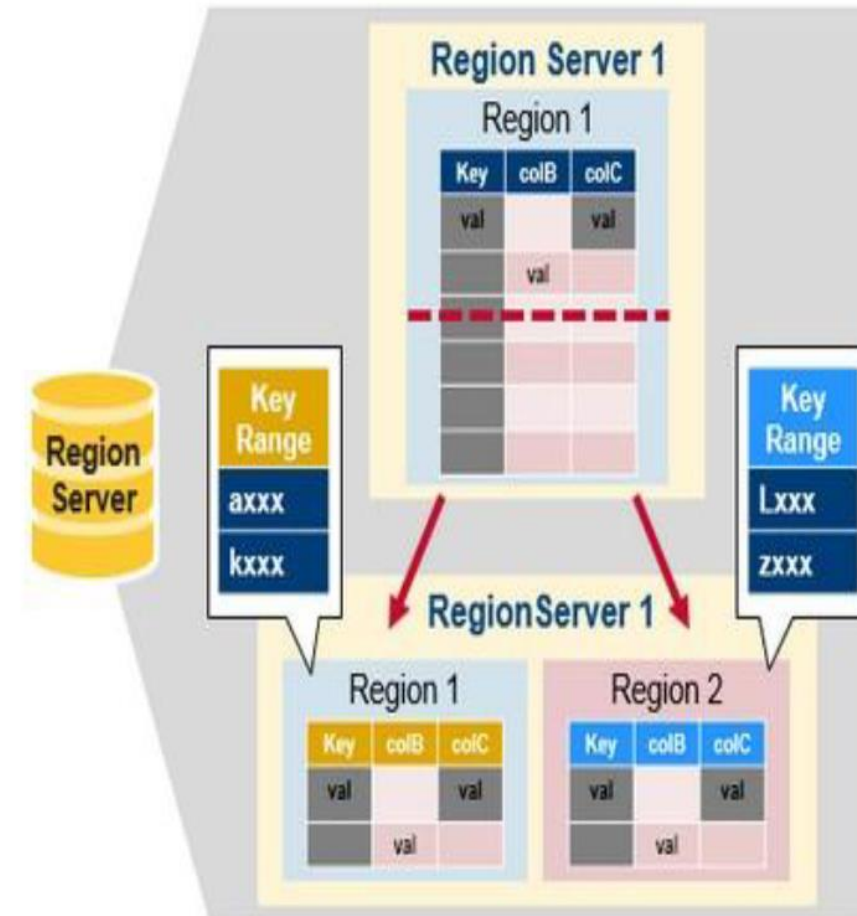
- They store the rows as sorted KeyValues on disk.



REGION SPLIT

Initially there is one region per table.

- When a region grows too large, it splits into two child regions.
- Both child regions, representing one-half of the original region, are opened in parallel on the same Region server, and then the split is reported to the HMaster.
- For load balancing reasons, the HMaster may schedule for new regions to be moved off to other servers.



when region size >
hbase.hregion.max.
filesize → split

APACHE HBASE ARCHITECTURE BENEFITS

HBase provides the following benefits:

- Strong consistency model- When a write returns, all readers will see same value
- Scales automatically- Regions split when data grows too large- Uses HDFS to spread and replicate data
- Built-in recovery- Using Write Ahead Log (similar to journaling on file system)
- Integrated with Hadoop- MapReduce on HBase is straightforward

RDBMS	HBASE
RDBMS is row-oriented databases	HBase is a distributed, column-oriented data storage system
RDBMS tables have fixed-schema	Hbase tables do not have fixed- schema
RDBMS tables guarantee ACID properties	Hbase tables guarantee consistency and partition tolerance
RDBMS uses SQL (Structured query Language) to query the data	Hbase uses Java client API and Jruby